

CUDA

```
cat>> ass1.cu
```

```
Nvcc ass1.cu
```

```
./a.out
```

1

MATRIX MULTIPLICATION

```
#include <iostream>
```

```
using namespace std;
```

```
// CUDA code to multiply matrices
```

```
__global__ void multiply(int* A, int* B, int* C, int size) {
```

```
    // Uses thread indices and block indices to compute each element
```

```
    int row = blockIdx.y * blockDim.y + threadIdx.y;
```

```
    int col = blockIdx.x * blockDim.x + threadIdx.x;
```

```
    if (row < size && col < size) {
```

```
        int sum = 0;
```

```
        for (int i = 0; i < size; i++) {
```

```
            sum += A[row * size + i] * B[i * size + col];
```

```
        }
```

```
        C[row * size + col] = sum;
```

```
    }
```

```
}
```

```
void initialize(int* matrix, int size) {  
    for (int i = 0; i < size * size; i++) {  
        matrix[i] = rand() % 10;  
    }  
}
```

```
void print(int* matrix, int size) {  
    for (int row = 0; row < size; row++) {  
        for (int col = 0; col < size; col++) {  
            cout << matrix[row * size + col] << " ";  
        }  
        cout << '\n';  
    }  
    cout << '\n';  
}
```

```
int main() {  
    int* A, * B, * C;
```

```
int N = 2;
int blockSize = 16;

int matrixSize = N * N;
size_t matrixBytes = matrixSize * sizeof(int);

A = new int[matrixSize];
B = new int[matrixSize];
C = new int[matrixSize];

initialize(A, N);
initialize(B, N);
cout << "Matrix A: \n";
print(A, N);

cout << "Matrix B: \n";
print(B, N);

int* X, * Y, * Z;
// Allocate space
cudaMalloc(&X, matrixBytes);
cudaMalloc(&Y, matrixBytes);
cudaMalloc(&Z, matrixBytes);
```

```
// Copy values from A to X
cudaMemcpy(X, A, matrixBytes, cudaMemcpyHostToDevice);

// Copy values from A to X and B to Y
cudaMemcpy(Y, B, matrixBytes, cudaMemcpyHostToDevice);

// Threads per CTA dimension
int THREADS = 2;

// Blocks per grid dimension (assumes THREADS divides N evenly)
int BLOCKS = N / THREADS;

// Use dim3 structs for block and grid dimensions
dim3 threads(THREADS, THREADS);
dim3 blocks(BLOCKS, BLOCKS);

// Launch kernel
multiply<<<blocks, threads>>>(X, Y, Z, N);

cudaMemcpy(C, Z, matrixBytes, cudaMemcpyDeviceToHost);
cout << "Multiplication of matrix A and B: \n";
print(C, N);
```

```
delete[] A;
delete[] B;
delete[] C;

cudaFree(X);
cudaFree(Y);
cudaFree(Z);

return 0;
}
```

VECTOR ADDITION

```
#include <iostream>
using namespace std;

__global__ void add(int* A, int* B, int* C, int size) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;

    if (tid < size) {
        C[tid] = A[tid] + B[tid];
    }
}
```

```
void initialize(int* vector, int size) {  
    for (int i = 0; i < size; i++) {  
        vector[i] = rand() % 10;  
    }  
}
```

```
void print(int* vector, int size) {  
    for (int i = 0; i < size; i++) {  
        cout << vector[i] << " ";  
    }  
    cout << endl;  
}
```

```
int main() {  
    int N = 4;  
    int* A, * B, * C;  
  
    int vectorSize = N;  
    size_t vectorBytes = vectorSize * sizeof(int);  
  
    A = new int[vectorSize];  
    B = new int[vectorSize];
```

```
C = new int[vectorSize];
```

```
initialize(A, vectorSize);
```

```
initialize(B, vectorSize);
```

```
cout << "Vector A: ";
```

```
print(A, N);
```

```
cout << "Vector B: ";
```

```
print(B, N);
```

```
int* X, * Y, * Z;
```

```
cudaMalloc(&X, vectorBytes);
```

```
cudaMalloc(&Y, vectorBytes);
```

```
cudaMalloc(&Z, vectorBytes);
```

```
cudaMemcpy(X, A, vectorBytes, cudaMemcpyHostToDevice);
```

```
cudaMemcpy(Y, B, vectorBytes, cudaMemcpyHostToDevice);
```

```
int threadsPerBlock = 256;
```

```
int blocksPerGrid = (N + threadsPerBlock - 1) / threadsPerBlock;
```

```
add<<<blocksPerGrid, threadsPerBlock>>>(X, Y, Z, N);
```

```
cudaMemcpy(C, Z, vectorBytes, cudaMemcpyDeviceToHost);
```

```
cout << "Addition: ";
```

```
print(C, N);
```

```
delete[] A;
```

```
delete[] B;
```

```
delete[] C;
```

```
cudaFree(X);
```

```
cudaFree(Y);
```

```
cudaFree(Z);
```

```
return 0;
```

```
}
```